

图 9.1 两种视觉神经元对图像信号的接收方式



图 9.2 大脑获得完整图像信息的方式




图 9.3 卷积神经网络示意图




图 9.4 滤镜示意图




图 9.5 计算卷积  $U * W$



图 9.6 颐和园佛香阁




图 9.7 佛香阁的横向轮廓




图 9.8 佛香阁的纵向轮廓






图 9.9 滤镜过滤  $U^{(1)}, U^{(2)}, \dots, U^{(l)}$  后镜像的叠加




图 9.10 当前层输出的镜像  $C^{(1)}, C^{(2)}, \dots, C^{(k)}$




图 9.11 最大值池化过程




图 9.12 最大值池化后佛香阁的横向边界轮廓




图 9.13 平均值池化后佛香阁的横向边界轮廓




图 9.14 卷积神经网络的结构

```

1  import tensorflow as tf
2  import numpy as np
3  import matplotlib.pyplot as plt
4  from sklearn.datasets import load_sample_image
5
6  china = load_sample_image("china.jpg")
7  height, width, channels = china.shape
8  images = china.reshape(1, height, width, channels)
9  plt.figure(9.7)
10 plt.imshow(china)
11
12 X = tf.placeholder(tf.float32, shape = [1, height, width, channels])
13 filters = np.zeros(shape = (2, 1, 3, 1), dtype = np.float32)
14 filters[0, :, :, :] = -1
15 filters[1, :, :, :] = 1
16 convolution = tf.nn.conv2d(X, filters, strides = [1,1,1,1], padding = 'VALID')
17 max_pool = tf.nn.max_pool(convolution, ksize = [1,10,10,1],
18                             strides = [1,10,10,1], padding = 'SAME')
19
20 with tf.Session() as sess:
21     conv_output = sess.run(convolution, feed_dict = {X : images})
22     pool_output = sess.run(max_pool, feed_dict = {X : images})
23     plt.figure(9.8)
24     plt.imshow(conv_output[0, :, :, 0], cmap = "gray")
25     plt.figure(9.12)
26     plt.imshow(pool_output[0, :, :, 0], cmap = "gray")

```

图 9.15 卷积层与池化




图 9.16 用填充全 0 的方法计算卷积





图 9.17 训练数据集中的图片采样

```
machine_learning.cnn.cats_and_dogs.data_reader

1  import numpy as np
2  import cv2
3  import os
4  import glob
5
6  def load_data(data_path, image_size, classes):
7      images = []
8      labels = []
9      for c in classes:
10         index = classes.index(c)
11         path = os.path.join(data_path, c+'.jpg')
12         files = glob.glob(path)
13         for file in files:
14             image = cv2.imread(file)
15             image = cv2.resize(image, (image_size, image_size), 0, 0, cv2.INTER_LINEAR)
16             image = np.multiply(image, 1.0 / 255.0)
17             images.append(image)
18             label = np.zeros(len(classes))
19             label[index] = 1
20             labels.append(label)
21     return np.array(images), np.array(labels)
```

图 9.18 读入猫和狗的图片

```

1  import tensorflow as tf
2  from sklearn.model_selection import train_test_split
3  from machine_learning.cnn.cats_and_dogs.data_reader import load_data
4
5  n_classes = 2
6  img_size = 128
7  n_channels = 3
8  X = tf.placeholder(tf.float32, shape = [None, img_size, img_size, n_channels])
9  y = tf.placeholder(tf.float32, shape = [None, n_classes])
10 conv1 = tf.layers.conv2d(X, filters = 32, kernel_size = [3,3],
11                           strides = [1,1], padding = 'SAME')
12 conv1_pool = tf.nn.max_pool(conv1, ksize = [1,2,2,1],
13                              strides = [1,2,2,1], padding = 'SAME')
14 conv2 = tf.layers.conv2d(conv1_pool, filters = 32, kernel_size = [3, 3],
15                           strides = [1,1], padding = 'SAME')
16 conv2_pool = tf.nn.max_pool(conv2, ksize = [1,2,2,1],
17                              strides = [1,2,2,1], padding = 'SAME')
18 conv3 = tf.layers.conv2d(conv2_pool, filters = 64, kernel_size = [3,3],
19                           strides=[1,1], padding = 'SAME')
20 conv3_pool = tf.nn.max_pool(conv3, ksize = [1,2,2,1],
21                              strides = [1,2,2,1], padding = 'SAME')
22 n_features = conv3_pool.get_shape()[1:4].num_elements()
23 conv_flat = tf.reshape(conv3_pool, [-1, n_features])
24 fc1 = tf.layers.dense(conv_flat, 128, activation=tf.nn.relu)
25 fc2 = tf.layers.dense(fc1, n_classes)
26 cross_entropy = tf.nn.softmax_cross_entropy_with_logits(logits=fc2, labels=y)
27 cost = tf.reduce_mean(cross_entropy)
28 optimizer = tf.train.AdamOptimizer(learning_rate=1e-4).minimize(cost)
29 y_pred = tf.nn.softmax(fc2)
30 correct_prediction = tf.equal(tf.argmax(y_pred, axis = 1), tf.argmax(y, axis = 1))
31 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
32
33 saver = tf.train.Saver()
34 with tf.Session() as sess:
35     tf.global_variables_initializer().run()
36     data_path = './ cats_and_dogs'


```

```


37 images, labels = load_data(data_path, img_size, ['dog','cat'])
38 image_train, image_test, label_train, label_test = train_test_split(
39     images, labels, test_size=0.1, random_state=0)
40 n_epochs = 20
41 batch_size = 32
42 best_accuracy = 0
43 for epoch in range(n_epochs):
44     for batch in range(len(image_train) // batch_size):
45         start = batch * batch_size
46         end = (batch + 1) * batch_size
47         X_batch = image_train[start:end]
48         y_batch = label_train[start:end]
49         sess.run(optimizer, feed_dict = {X:X_batch, y:y_batch})
50     test_accuracy = accuracy.eval(feed_dict = {X:image_test, y:label_test})
51     print("epoch {}: test accuracy = {}".format(epoch, test_accuracy))
52     if test_accuracy > best_accuracy:
53         saver.save(sess, "./cats_and_dogs_model/cats_and_dogs.ckpt")
54         best_accuracy = test_accuracy

```

图 9.19 猫和狗图片识别问题的卷积神经网络




$p = (0.72, 0.28)$   
(a)



$p = (0.01, 0.99)$   
(b)



$p = (0.47, 0.53)$   
(c)



$p = (0.99, 0.01)$   
(d)

图 9.20 卷积神经网络对 4 个测试数据的概率预测

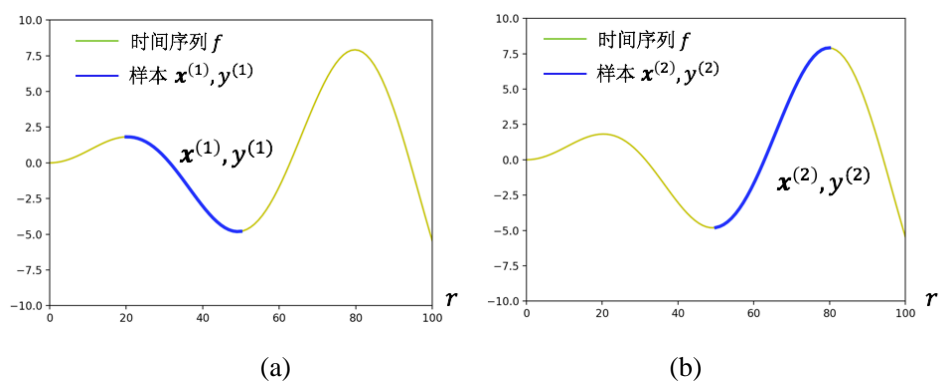


图 9.21 起点是 20 和 50 的两个长度为 30 的时间序列




图 9.22 记忆单元的示意图

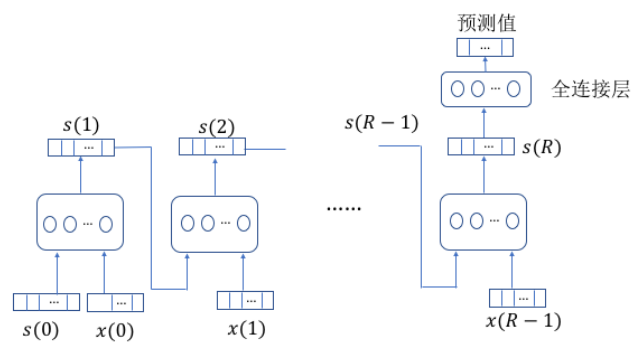


图 9.23 循环神经网络的结构示意图






图 9.24 多层循环神经网络的结构




图 9.25 多层循环神经网络的第  $l$  层示意图

```

1  import tensorflow as tf
2  import numpy as np
3
4  def time_series(r):
5      return r / 10.0 * np.sin(r / 10.0)
6
7  def get_samples(n_samples, r_max, R):
8      r0 = np.random.rand(n_samples, 1) * (r_max - R)
9      r = r0 + np.arange(0, R + 1)
10     f = time_series(r)
11     x = f[:, 0:R].reshape(-1, R, 1)
12     y = f[:, R].reshape(-1, 1)
13     return x, y
14
15     n_inputs = 1
16     n_outputs = 1
17     X = tf.placeholder(tf.float32, [None, 30, n_inputs])
18     y = tf.placeholder(tf.float32, [None, n_outputs])
19     cell = tf.contrib.rnn.BasicRNNCell(num_units = 50, activation = tf.nn.relu)
20     states, final_state = tf.nn.dynamic_rnn(cell, X, dtype = tf.float32)
21     preds = tf.layers.dense(final_state, n_outputs)
22     loss = tf.reduce_mean(tf.square(preds - y))
23     optimizer = tf.train.AdamOptimizer(learning_rate = 0.001)
24     training_op = optimizer.minimize(loss)
25
26     with tf.Session() as sess:
27         tf.global_variables_initializer().run()
28         R, rmax = 30, 100
29         X_train, y_train = get_samples(100, r_max, R)
30         X_test, y_test = get_samples(100, r_max, R)
31         for iteration in range(1000):
32             sess.run(training_op, feed_dict = {X: X_train, y: y_train})
33             if iteration % 100 == 0:
34                 mse = loss.eval(feed_dict = {X: X_test, y: y_test})
35                 print("iteration {} : MSE = {}".format(iteration, mse))

```

图 9.26 时间序列自回归问题的循环神经网络算法




图 9.27 样本时间序列标签值与预测值

```

1  import tensorflow as tf
2  import numpy as np
3
4  def time_series(r):
5      return r / 10.0 * np.sin(r / 10.0)
6
7  def get_samples(n_samples, r_max, R):
8      r0 = np.random.rand(n_samples, 1) * (r_max - R)
9      r = r0 + np.arange(0, R + 1)
10     f = time_series(r)
11     x = f[:, 0:R].reshape(-1, R, 1)
12     y = f[:, R].reshape(-1, 1)
13     return x, y
14
15  n_inputs = 1
16  n_outputs = 1
17  X = tf.placeholder(tf.float32, [None, R, n_inputs])
18  y = tf.placeholder(tf.float32, [None, n_outputs])
19  layers = []
20  layer_1 = tf.contrib.rnn.BasicRNNCell(num_units = 50, activation = tf.nn.relu)
21  layer_2 = tf.contrib.rnn.BasicRNNCell(num_units = 20, activation = tf.nn.relu)
22  layers.append(layer_1)
23  layers.append(layer_2)
24  multi_layer_cell = tf.contrib.rnn.MultiRNNCell(layers)
25  states, final_state = tf.nn.dynamic_rnn(multi_layer_cell, X, dtype = tf.float32)
26  preds = tf.layers.dense(final_state[-1], n_outputs)
27  loss = tf.reduce_mean(tf.square(preds - y))
28  optimizer = tf.train.AdamOptimizer(learning_rate = 0.001)
29  training_op = optimizer.minimize(loss)
30
31  with tf.Session() as sess:
32      tf.global_variables_initializer().run()
33      R, r_max = 30, 100
34      X_train, y_train = get_samples(100, r_max, R)
35      X_test, y_test = get_samples(100, r_max, R)
36      for iteration in range(1000):
37          sess.run(training_op, feed_dict={X: X_train, y: y_train})
38          if iteration % 100 == 0:
39              mse = loss.eval(feed_dict={X: X_test, y: y_test})
40              print("iteration { } : MSE = {}".format(iteration, mse))

```

图 9.28 时间序列自回归问题的多层循环神经网络算法




图 9.29 循环神经网络中的反向传播算法

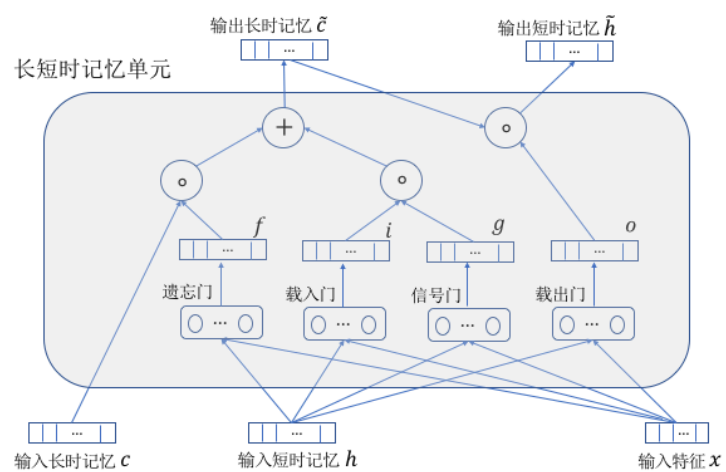


图 9.30 长短时记忆单元的内部结构

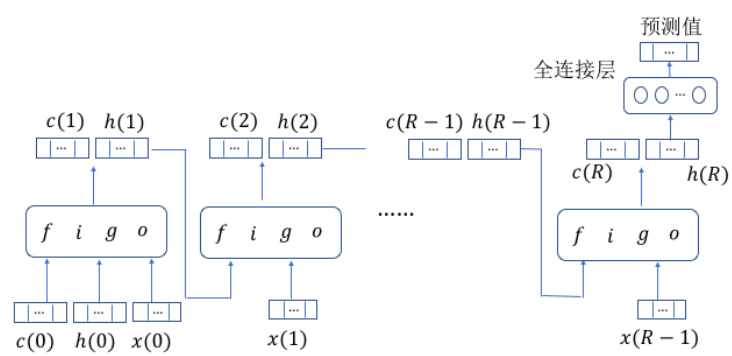


图 9.31 含有长短时记忆单元的循环神经网络示意图



```

1  import tensorflow as tf
2  import numpy as np
3
4  def time_series(r):
5      return r / 10.0 * np.sin(r / 10.0)
6
7  def get_samples(n_samples, r_max, R):
8      r0 = np.random.rand(n_samples, 1) * (r_max - R)
9      r = r0 + np.arange(0.0, R + 1)
10     f = time_series(r)
11     x = f[:, 0:R].reshape(-1, R, 1)
12     y = f[:, R].reshape(-1, 1)
13     return x, y
14
15     n_inputs = 1
16     n_outputs = 1
17     num_units = 50
18     X = tf.placeholder(tf.float32, [None, R, n_inputs])
19     y = tf.placeholder(tf.float32, [None, n_outputs])
20     lstm_cell = tf.contrib.rnn.BasicLSTMCell(num_units)
21     states, final_state = tf.nn.dynamic_rnn(lstm_cell, X, dtype = tf.float32)
22     preds = tf.layers.dense(final_state.h, n_outputs)
23     loss = tf.reduce_mean(tf.square(preds - y))
24     optimizer = tf.train.AdamOptimizer(learning_rate = 0.001)
25     training_op = optimizer.minimize(loss)
26
27     with tf.Session() as sess:
28         tf.global_variables_initializer().run()
29         R, r_max = 30, 100
30         X_train, y_train = get_samples(100, r_max, R)
31         X_test, y_test = get_samples(100, r_max, R)
32         for iteration in range(1000):
33             sess.run(training_op, feed_dict={X: X_train, y: y_train})
34             if iteration % 100 == 0:
35                 mse = loss.eval(feed_dict={X: X_test, y: y_test})
36                 print("iteration { } : MSE = {}".format(iteration, mse))

```

图 9.32 含有长短时记忆单元的循环神经网络算法